


```
# data analysis
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")

# read data
df=pd.read_csv("weatherAUS.csv")
df
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGus
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	\
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	\
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	
...	
140580	2012-08-08	Darwin	17.8	31.4	0.0	7.0	11.1	
140581	2012-08-09	Darwin	18.4	32.2	0.0	7.4	11.1	
140582	2012-08-10	Darwin	17.3	32.8	0.0	7.6	10.9	
140583	2012-08-11	Darwin	17.7	31.6	0.0	7.4	11.1	
140584	2012-08-08	NaN	NaN	NaN	NaN	NaN	NaN	

140585 rows × 23 columns



```
# data information
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 140585 entries, 0 to 140584
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   140585 non-null object
1   Location               140584 non-null object
2   MinTemp                139184 non-null float64
3   MaxTemp                139371 non-null float64
4   Rainfall              137397 non-null float64
5   Evaporation            79490 non-null float64
6   Sunshine               73907 non-null float64
7   WindGustDir            130377 non-null object
8   WindGustSpeed          130436 non-null float64
9   WindDir9am             130137 non-null object
10  WindDir3pm              136409 non-null object
11  WindSpeed9am            138855 non-null float64
12  WindSpeed3pm            137559 non-null float64
13  Humidity9am             138026 non-null float64
14  Humidity3pm             136918 non-null float64
15  Pressure9am             125530 non-null float64
16  Pressure3pm             125566 non-null float64
17  Cloud9am                85968 non-null float64
18  Cloud3pm                83112 non-null float64
19  Temp9am                 138874 non-null float64
20  Temp3pm                 137704 non-null float64
21  RainToday               137397 non-null object
22  RainTomorrow            137394 non-null object
dtypes: float64(16), object(7)
memory usage: 24.7+ MB
```

```
# data contain float64(16), object(7) and 140585 entries are available in dataset
```

```
# check null values available or not
df.isnull().sum()
```

```

Date          0
Location       1
MinTemp       1401
MaxTemp       1214
Rainfall      3188
Evaporation   61095
Sunshine      66678
WindGustDir    10208
WindGustSpeed  10149
WindDir9am     10448
WindDir3pm     4176
WindSpeed9am   1730
WindSpeed3pm   3026
Humidity9am    2559
Humidity3pm    3667
Pressure9am    15055
Pressure3pm    15019
Cloud9am       54617
Cloud3pm       57473
Temp9am        1711
Temp3pm        2881
RainToday      3188
RainTomorrow   3191
dtype: int64
```

```
# data shows that so many null values are available in dataset
```

```
# handling missing values
from sklearn.impute import SimpleImputer
si= SimpleImputer(missing_values=np.nan, strategy='mean')
```

```
num_col=df.select_dtypes(["float","int"]).columns
```

```
df[num_col]=si.fit_transform(df[num_col])
```

```
cat_col=df.select_dtypes(["O"]).columns
```

```
si= SimpleImputer(missing_values=np.nan, strategy='most_frequent')
df[cat_col]=si.fit_transform(df[cat_col])
```

```
df.head()
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir
0	2008-12-01	Albury	13.4	22.9	0.6	5.413627	7.585605	W
1	2008-12-02	Albury	7.4	25.1	0.0	5.413627	7.585605	WNW
2	2008-12-03	Albury	12.9	25.7	0.0	5.413627	7.585605	WSW
3	2008-12-04	Albury	9.2	28.0	0.0	5.413627	7.585605	NE
4	2008-12-05	Albury	17.5	32.3	1.0	5.413627	7.585605	W

5 rows × 23 columns

```
df.isnull().sum()
```

```

Date          0
Location       0
MinTemp       0
MaxTemp       0
Rainfall      0
Evaporation   0
Sunshine      0
WindGustDir    0
WindGustSpeed  0
WindDir9am     0
WindDir3pm     0
WindSpeed9am   0
WindSpeed3pm   0
Humidity9am    0
```

```

Humidity3pm      0
Pressure9am      0
Pressure3pm      0
Cloud9am         0
Cloud3pm         0
Temp9am          0
Temp3pm          0
RainToday        0
RainTomorrow     0
dtype: int64

# convert object columns in numerical form
from sklearn.preprocessing import LabelEncoder
lr=LabelEncoder()

cat_col

Index(['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm',
      'RainToday', 'RainTomorrow'],
      dtype='object')

for i in cat_col:
    df[i]=lr.fit_transform(df[i])

# splitting data into train_test_split
x=df.iloc[:, :-1].values
x

array([[ 396.      ,  2.      , 13.4      , ..., 16.9      ,
        21.8      ,  0.      ],
       [ 397.      ,  2.      ,  7.4      , ..., 17.2      ,
        24.3      ,  0.      ],
       [ 398.      ,  2.      , 12.9      , ..., 21.      ,
        23.2      ,  0.      ],
       ...,
       [1715.     , 13.      , 17.3      , ..., 24.1      ,
        32.      ,  0.      ],
       [1716.     , 13.      , 17.7      , ..., 21.7      ,
        30.4      ,  0.      ],
       [1705.     ,  9.      , 11.93831978, ..., 16.70854588,
        21.40531865,  0.      ]])

y=df.iloc[:, -1].values
y

array([0, 0, 0, ..., 0, 0, 0])

# scaling
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x=sc.fit_transform(x)

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=145)

# creating a model
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense

# using early stoping concept
from tensorflow.keras.callbacks import EarlyStopping

early_stop=EarlyStopping(monitor='val_loss',mode='min',verbose=1,patience=25)

# initialize the model
ann=Sequential()

# add layers to the model
ann.add(Dense(60,activation='relu'))
ann.add(Dense(60,activation='relu'))
# output layer
ann.add(Dense(1,activation='sigmoid'))

# established the connection between the layers

```

```
ann.compile(optimizer='sgd',loss='binary_crossentropy',metrics=['accuracy'])
```

```
# train the model
```

```
ann.fit(x_train,y_train,epochs=200,validation_data=(x_test,y_test),batch_size=125,callbacks=[early_stop])
```

```
Epoch 7/200
900/900 [=====] - 2s 2ms/step - loss: 0.2979 - accuracy: 0.8730 - val_loss: 0.3578 - val_accuracy: 0.848
Epoch 8/200
900/900 [=====] - 2s 2ms/step - loss: 0.2979 - accuracy: 0.8730 - val_loss: 0.3578 - val_accuracy: 0.848
Epoch 9/200
900/900 [=====] - 2s 2ms/step - loss: 0.2979 - accuracy: 0.8730 - val_loss: 0.3577 - val_accuracy: 0.848
Epoch 10/200
900/900 [=====] - 2s 3ms/step - loss: 0.2979 - accuracy: 0.8729 - val_loss: 0.3577 - val_accuracy: 0.848
Epoch 11/200
900/900 [=====] - 2s 3ms/step - loss: 0.2979 - accuracy: 0.8727 - val_loss: 0.3577 - val_accuracy: 0.848
Epoch 12/200
900/900 [=====] - 2s 2ms/step - loss: 0.2979 - accuracy: 0.8730 - val_loss: 0.3578 - val_accuracy: 0.848
Epoch 13/200
900/900 [=====] - 2s 2ms/step - loss: 0.2978 - accuracy: 0.8730 - val_loss: 0.3580 - val_accuracy: 0.848
Epoch 14/200
900/900 [=====] - 2s 2ms/step - loss: 0.2978 - accuracy: 0.8727 - val_loss: 0.3578 - val_accuracy: 0.848
Epoch 15/200
900/900 [=====] - 2s 2ms/step - loss: 0.2978 - accuracy: 0.8726 - val_loss: 0.3578 - val_accuracy: 0.848
Epoch 16/200
900/900 [=====] - 2s 3ms/step - loss: 0.2978 - accuracy: 0.8729 - val_loss: 0.3580 - val_accuracy: 0.848
Epoch 17/200
900/900 [=====] - 2s 3ms/step - loss: 0.2978 - accuracy: 0.8728 - val_loss: 0.3579 - val_accuracy: 0.847
Epoch 18/200
900/900 [=====] - 2s 2ms/step - loss: 0.2978 - accuracy: 0.8729 - val_loss: 0.3579 - val_accuracy: 0.848
Epoch 19/200
900/900 [=====] - 2s 2ms/step - loss: 0.2977 - accuracy: 0.8726 - val_loss: 0.3579 - val_accuracy: 0.847
Epoch 20/200
900/900 [=====] - 2s 2ms/step - loss: 0.2977 - accuracy: 0.8729 - val_loss: 0.3579 - val_accuracy: 0.848
Epoch 21/200
900/900 [=====] - 2s 2ms/step - loss: 0.2977 - accuracy: 0.8731 - val_loss: 0.3580 - val_accuracy: 0.848
Epoch 22/200
900/900 [=====] - 2s 2ms/step - loss: 0.2977 - accuracy: 0.8729 - val_loss: 0.3580 - val_accuracy: 0.848
Epoch 23/200
900/900 [=====] - 2s 3ms/step - loss: 0.2977 - accuracy: 0.8732 - val_loss: 0.3582 - val_accuracy: 0.848
Epoch 24/200
900/900 [=====] - 2s 2ms/step - loss: 0.2977 - accuracy: 0.8729 - val_loss: 0.3581 - val_accuracy: 0.848
Epoch 25/200
900/900 [=====] - 2s 2ms/step - loss: 0.2977 - accuracy: 0.8731 - val_loss: 0.3582 - val_accuracy: 0.847
Epoch 26/200
900/900 [=====] - 2s 2ms/step - loss: 0.2976 - accuracy: 0.8729 - val_loss: 0.3581 - val_accuracy: 0.848
Epoch 27/200
900/900 [=====] - 2s 2ms/step - loss: 0.2976 - accuracy: 0.8727 - val_loss: 0.3580 - val_accuracy: 0.847
Epoch 28/200
900/900 [=====] - 2s 2ms/step - loss: 0.2976 - accuracy: 0.8731 - val_loss: 0.3582 - val_accuracy: 0.847
Epoch 29/200
900/900 [=====] - 3s 3ms/step - loss: 0.2976 - accuracy: 0.8728 - val_loss: 0.3582 - val_accuracy: 0.848
Epoch 30/200
900/900 [=====] - 2s 2ms/step - loss: 0.2975 - accuracy: 0.8730 - val_loss: 0.3584 - val_accuracy: 0.848
Epoch 31/200
900/900 [=====] - 2s 2ms/step - loss: 0.2975 - accuracy: 0.8728 - val_loss: 0.3582 - val_accuracy: 0.847
Epoch 32/200
900/900 [=====] - 2s 2ms/step - loss: 0.2975 - accuracy: 0.8729 - val_loss: 0.3585 - val_accuracy: 0.847
Epoch 33/200
900/900 [=====] - 2s 2ms/step - loss: 0.2975 - accuracy: 0.8730 - val_loss: 0.3584 - val_accuracy: 0.848
Epoch 34/200
900/900 [=====] - 2s 2ms/step - loss: 0.2975 - accuracy: 0.8727 - val_loss: 0.3582 - val_accuracy: 0.847
Epoch 34: early stopping
keras.callbacks.History at 0x7f27f4049a0>
```

```
ann.history.history
```

```
0.558199903010090593},
'val_accuracy': [0.8496994972229004,
0.8493082523345947,
0.8483124375343323,
0.8493082523345947,
0.8482056856155396,
0.8484191298484802,
0.8486680388450623,
0.8489170074462891,
0.8483479619026184,
0.8485613465309143,
0.8488814830780029,
0.8486325144767761,
0.848845899105072,
0.8486325144767761,
0.8485613465309143,
0.8484546542167664,
0.8475655317306519,
0.8485613465309143,
0.8478856086730957,
0.8485258221626282,
0.8481701612472534,
0.8480989933013916,
0.8484902381896973,
0.848774790763855,
0.8479211926460266,
0.8486325144767761,
0.8477789163589478,
0.8478500843048096,
0.8481701612472534,
0.8488814830780029,
0.8474944233894348,
0.8477078080177307,
0.848845899105072,
0.8472809791564941]]}
```

```
ypred=ann.predict(x_test)
ypred=ypred>0.5

879/879 [=====] - 1s 1ms/step
```

```
# evaluation on test data
from sklearn.metrics import classification_report
print(classification_report(ypred,y_test))
```

	precision	recall	f1-score	support
False	0.93	0.88	0.90	23224
True	0.55	0.70	0.61	4893
accuracy			0.85	28117
macro avg	0.74	0.79	0.76	28117
weighted avg	0.87	0.85	0.85	28117

```
# model predicted 85% accuracy.
```